1 Result

This section presents the findings of our user study, evaluating the performance of an AI-driven natural language interface (NLI) against a conventional search system in an e-commerce context. The results address user satisfaction, the role of filters, technical feasibility, and possible correction mechanisms for query optimization, providing insights into the effectiveness of dynamic SQL query generation compared to static templates.

1.1 User Satisfaction and Search Behavior(RSQ1)

Figure **??** illustrates that all participants rated the AI-driven system's results highly on a 5-point Likert scale, with 69% (9/13) selecting 'very satisfied' (1/5) and 31% (4/13) 'satisfied' (2/5), resulting in a mean of 1.69 (SD 0.48, Bessel-corrected). This low variability reflects consistent user approval, attributed to the system's accurate intent interpretation, including its handling of synonyms (e.g., 'Wiener Dog' as 'Dachshund,' Fig. **??**). This feature was absent in the conventional system. In contrast, satisfaction with the conventional system was markedly lower, averaging 3.62(SD = 0.51) on the same scale,

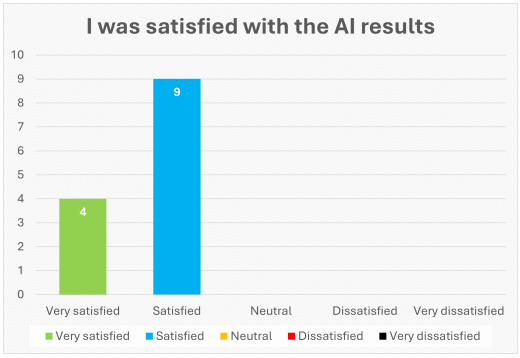

Fig. 1. The users' rating regarding the satisfaction of the results displayed by the AI. As can be seen in the picture, the participants were very satisfied with the results of the AI.

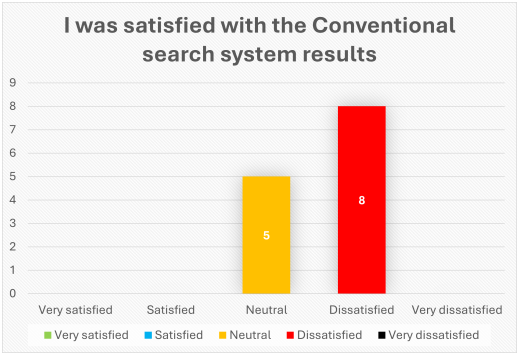

Fig. 2. The evaluation of the users regarding the satisfaction with the result displayed by the conventional system. The picture shows that they were not satisfied with the results displayed by the conventional system.

with only 38% (5 out of 13) rating it as neutral (3/5), and 62% (7 out of 13) as dissatisfied (4/5), with no participants rating it as very high,high or very dissatisfied (see Figure **??**). Users found its performance lacking, particularly when compared to the AI's intuitive query handling. An interesting observation emerged: as participants grew accustomed to the AI's ability to understand their intent, their reliance on manual filters decreased. This shift likely contributed to the conventional system's perceived inferiority, as it depended heavily on filter usage to refine results. This rapid adaptation suggests a potential change in search behavior, though further analysis is needed to confirm its extent. It's worth noting that the results for the conventional system should be considered in light of the relatively simple static template we used. This basic design may have contributed to its poorer performance, and a more sophisticated static template could potentially have yielded better outcomes.

1.2 The Role of Filters in Enhancing User Experience

While the AI demonstrates potential to supplant traditional filters, our study underscores their enduring value. Figure **??** details preferences from a 5-point Likert scale (1='strongly agree', 5='strongly disagree'): 69.2% (9/13) agreed or

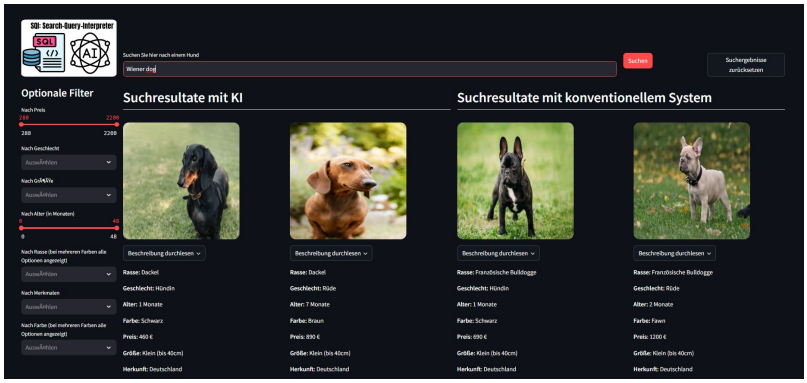

Fig. 3. In this picture we see the website that was developed for the project. The website was divided so that the user can quickly see which results belong to AI and which to the conventional system. To demonstrate the performance of the AI, Wiener dog was entered in the input field. As you can see, the AI can handle this, but the conventional system cannot.

strongly agreed they could forgo filters (scores 1-2), 23.1% (3/13) disagreed, favoring retention (score 4), and 7.7% (1/13) were neutral (score 3). Open-ended responses further suggest filters complement the AI for specific tasks (e.g., price sorting). Of the participants who preferred filters, three stated in open-ended comments that they appreciated the visual clarity, such as when sorting by price or selecting ranges (e.g., 'price between 50-100€'), compared to text-based inputs. For instance, participants noted that setting a price range via a slider was faster and more intuitive than specifying it textually. These findings, however, are based on a small, homogeneous sample of 13 participants, primarily computer science students, which may limit generalizability.

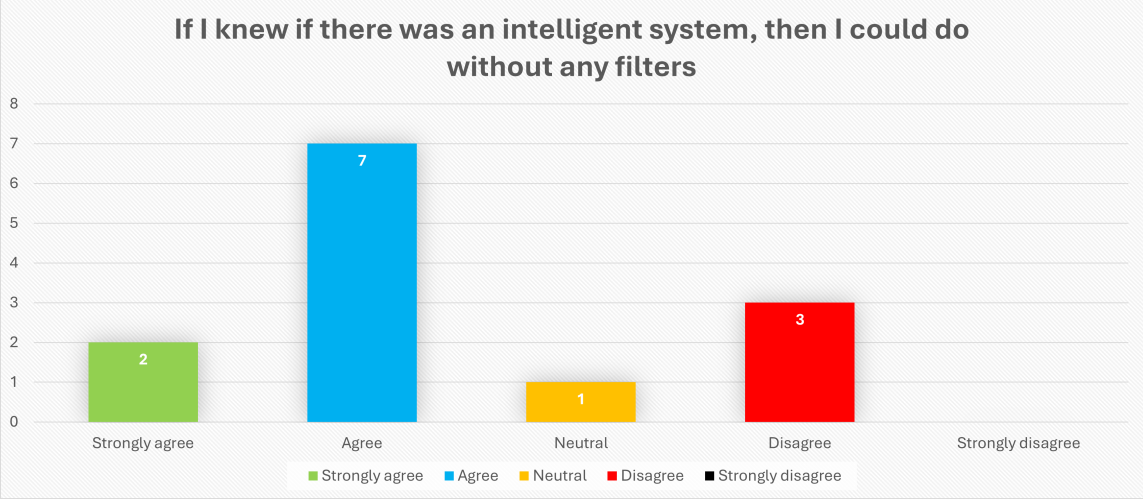

Fig. 4. After the users have used the system, they should state whether they would manage without the filter.The results show that the majority of users can manage without a filter without any problems. Three times "disagree" was indicated, while no "strongly disagree" was recorded.

1.3 Technical Feasibility and Security Considerations (RSQ2)

Building on user experience insights, this subsection examines the technical implications of the AI-driven system. During the study, some participants tested the system's boundaries, revealing critical security concerns. For example, one user entered a command to 'delete all tables,' which the AI translated into an executable SQL query—an unintended capability that exposed a vulnerability. Ideally, the system should restrict operations to read-only searches, preventing modifications to the database or its schema. While our simple dog webshop schema excluded sensitive data like user information, real-world applications would demand robust safeguards. To mitigate this vulnerability, we propose implementing a blacklist configuration to restrict the AI system's SQL query generation. This file (e.g. in csv-format) would specify tables or columns deemed off-limits, ensuring they remain inaccessible within the generated queries. Additionally, given that search results constitute a read-only operation, we recommend incorporating a list of prohibited keywords—such as 'DELETE,' 'UPDATE,' or 'DROP'—into the lock-file. Prior to execution, each AI-generated SQL query would be validated against these restricted elements (tables, columns, and keywords), for example through the use of regular expressions, to guarantee compliance with security and regulatory standards.

1.4 Optimizing Query Correction Strategies (RSQ3)

To address usability challenges and enhance trust, we investigated user preferences for refining inaccurate search results. Participants suggested several strategies, detailed below, to improve the system's transparency and responsiveness.

1.4.1 Dynamic Filter Adjustment. A participant has proposed an adjustment of the filters, based on the input provided. As illustrated in Figure **??**, the idea is to highlight implicitely used filters visually, i.e. dynamically adjusting filters on the interface. Therefore, the user is able to ascertain which filters the AI utilizes and, consequently, identify the potential origin of an error. Therefore, the user has the option of either utilize the filter to resolve the issue or adjusting the initial search query.

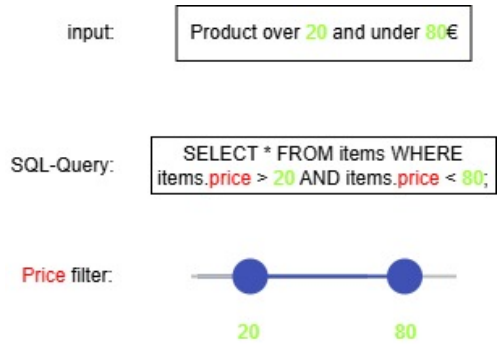

Fig. 5. Here the user can see how the filter is adapted to the user's input. The filter adapts to user input, revealing AI misinterpretations (implicitely).

1.4.3 Display of Similar Results. A further point to be considered is the potential for the AI to present analogous products in the event that the search yields few results. A relevant example would be a search for dogs that are of medium size. In the event that the available results are limited, the artificial intelligence could be programmed to display dogs of smaller stature.

1.4.5 Query Confidence Indicator. Another suggestion from users is that the AI should more clearly indicate its uncertainty when interpreting a query, for example, by providing an uncertainty score/confidence ranging from 0 to 1. This score would reflect how confident the AI is in its understanding of the request. This would allow users to better assess whether the provided results might be inaccurate or incorrect, enabling them to make adjustments if needed.

1.4.2 User Notification for No Results. A further potential avenue for enhancing the efficancy of the user's outcomes is the implementation of an artificial intelligence system that can evaluate the combinability of diverse features during the user's input phase. In such a case, it is essential that the user be alerted to this possibility. The current state of the AI, reflecting the original user input, is shown in Figure **??**. The subsequent version has been enhanced to alert the user to the absence of search results for a given combination of features. In this instance, the "under 2 months" feature is distinctly emphasized, as it does not result in any search results. Therefore, the user has the capacity to modify the preliminary search query and discern the elements that are not compatible.

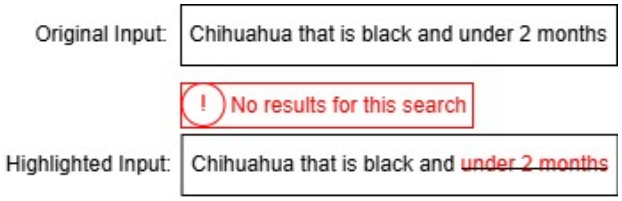

Fig. 6. Here the user is informed by the AI if the features are not combinable. It is made clear that 'under 2 months' cannot be combined with the other features, as this feature is crossed out.

1.4.4 Real-time AI-driven Suggestions. Another idea for improving the system is to offer the user search suggestions as they type. This not only enables the AI to deliver more relevant results, but also to better understand what the user is looking for. An example of this is shown in Figure **??**. The user enters 'Puppy that is black and...' and the AI suggests narrowing the search - for example to 'Puppy that is black and under three months'.

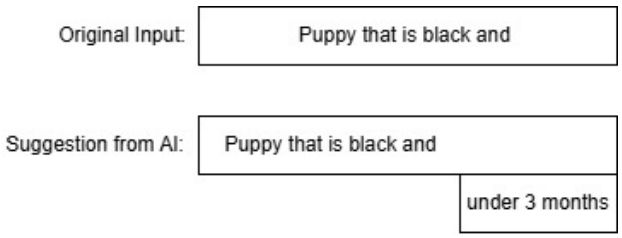

Fig. 7. The AI provides real-time suggestions for narrowing searches, such as additional features for 'Chihuahuas'.

An example from the current AI system illustrates this. When a user searches for a "cheap dog," the Claude API simply returns the least expensive dog available. However, this may not align with what the user intended; perhaps they were looking for a cost-effective yet suitable dog for specific needs. An uncertainty score could indicate how reliable the results are, signaling to the user that the interpretation of their query might not be entirely accurate. This would give the user the opportunity to refine their request for more appropriate responses.

Temporary page!

LaTeX was unable to guess the total number of pages correctly. As there was some unprocessed data that should have been added to the final page this extra page has been added to receive it.

If you rerun the document (without altering it) this surplus page will go away, because LaTeX now knows how many pages to expect for this document.